

# Time Series Modeling using Neural Networks

Dušan Fedorčák  
09/2022



# Background

- Ph.D. in computer science at VŠB-TU Ostrava
  - Neural networks & unsupervised self-organization
- Experienced in simulations
  - flood prediction system for MSK
  - traffic monitoring & prediction systems
- Experienced in computer graphics & scientific visualization
  - GIS related real-time 3D visualizations
- 5+ years in applied ML and artificial intelligence
  - Lead researcher in GoodAI – general artificial intelligence
  - CTO in Neuron Soundware – sound processing via Deep Learning
  - Lead ML in Merlon Intelligence Inc. – natural language processing

# Content

## DAY 1

- **Classical time series analysis**
  - *Decomposition of time series*
  - *ARIMA models family*
  - *State space models generalization*
- **Theoretical window**
  - *Neural Networks & Recurrent NNs*
  - *Time series specifics*
- **Practical examples**
  - *Simple regression – toy example*
  - *Rainfall-runoff simulation – regression*

– lunch break –

- **Practical examples**
  - *Trampoline jumping – classification*
  - *Local Weather Forecast – regression*


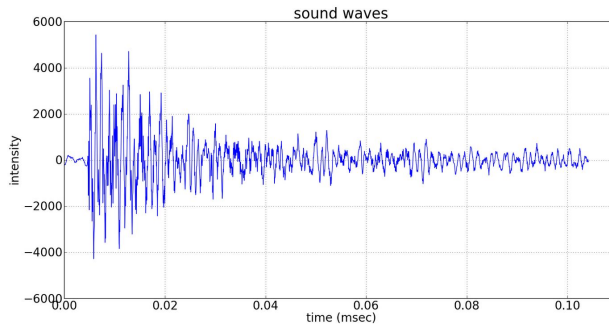

## DAY 2

- **Product Design & ML**
  - *Integration of ML models into products*
  - *Tips & tricks for debugging NNs*
- **Practical Examples** (in random order)
  - *Exoplanets Hunting*
  - *Mobile Motion Sensing*

– lunch break –

- *Manufacturing Process Modeling*
- *Financial distress prediction*
- [Google Drive Folder](#) with data
- [GitHub repository](#) with example sources
- [this presentation](#)










# Time series – example data



CandleStick Chart






# Time series – example tasks



# Time Series – classical analysis & modeling

- Time Series Decomposition
  - Inflation, trend, seasonality, differencing
- ARIMA models
  - <http://people.duke.edu/~rnau/411home.htm>



# Time Series – classical analysis & modeling

1. Open [this link](#)

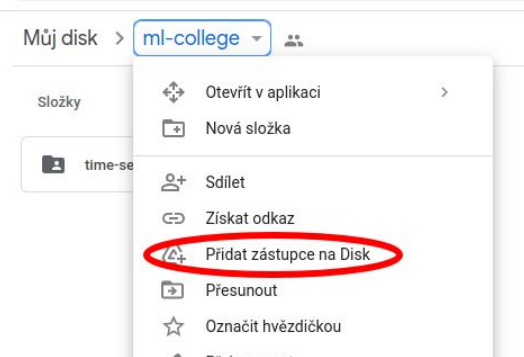
or

1. Open [\*colab.google.com\*](https://colab.google.com)

2. In the Open Notebook menu

- Navigate to GitHub tab
- Enter **mlcollege** into search bar
- Select **mlcollege/timeseriesanalysis** repo
- Open **arima-complete.ipynb**

*Do not forget to add this  
[Google Drive Folder](#) to your  
drive*




# State Space Models

- State Space Models
  - A dynamic system that **evolves over time**
  - Knowing **the current state of the model is enough** to predict the future
  - The true state of the system might **not** be **directly observable**
- Model Description
  - State
    - $\mathbf{x}_t \sim N(\mathbf{x}_t, \mathbf{P}_t)$
  - State Equation
    - $\mathbf{x}_t = \mathbf{F}\mathbf{x}_{t-1} + N(0, \mathbf{Q})$  – sometimes without noise
  - Observation Equation
    - $\mathbf{y}_t = \mathbf{H}\mathbf{x}_t + N(0, \mathbf{R})$



# Kalman Filter

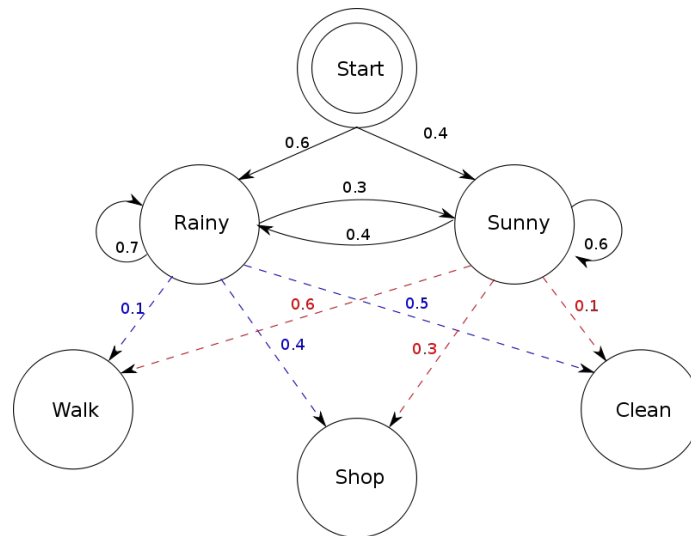
- Evolve state
  - $\mathbf{x}'_t = F\mathbf{x}_{t-1}$
  - $\mathbf{P}'_t = F\mathbf{P}_{t-1}F^T + Q$
- Integrate observation
  - $\mathbf{x}_t = \mathbf{x}'_t + K_t(y_t - H\mathbf{x}'_t)$
  - $\mathbf{P}_t = (I - K_tH)\mathbf{P}'_t$
- Kalman Gain
  - $K_t = \mathbf{P}'_t H^T (H\mathbf{P}'_t H^T + R)^{-1}$
- ARIMA and Kalman Filter
  - ARIMA can be viewed as a state space model
  - ARIMA can be fitted with MLE via Kalman Filter
  - <https://bookdown.org/rdpeng/timeseriesbook/maximum-likelihood-with-the-kalman-filter.html>
  - <https://towardsdatascience.com/the-kalman-filter-and-maximum-likelihood-9861666f6742>



# Hidden Markov Model

- Model Description

- HMM ( $\lambda$ ) can be viewed as a **state space model**
- Finite set of hidden states
  - $Q = \{q_1, q_2, \dots, q_n\}$ ,  $\pi = \{\pi_1, \pi_2, \dots, \pi_n\}$  – init
  - $n$  – number of states (hyperparameter)
- Set of observations
  - $O_i = (o^1, o^2, o^3, \dots, o^T)$
- Transition probability matrix & emissions
  - $A = (a_{00}, \dots, a_{nn})$ ,  $B = q_i \rightarrow o$



- Model Capabilities

- $P(O|\lambda)$  – Give prob. of  $O$  being produced by  $\lambda$  – *forward-backward alg.*
- $P(q_1, \dots, q_t|O, \lambda)$  – Give most likely sequence of states for given  $O$  – *Viterbi alg.*
- $O \Rightarrow \lambda$  – Model must be trainable with  $O$  – *Baum-Welch alg.*


# Kalman Filter vs. Hidden Markov Model

- **Kalman Filter**

- Continuous state
- Generic state & observation equation
- Linear dynamic system
- *Fusion of sensor readings and controls*
- *ARMA models implementation*

- **Hidden Markov Model**


- Discrete set of states
- N-states hyperparameter
- Emission & Transition tables
- *Speech recognition*
- *Time series segmentation*



# Time Series – goals in classical terminology



- **Forecasting**
  - Given the past and the present observation, what will the future look like?
- **Time scale analysis**
  - Given the observations, what time scales dominate when observing temporal variation in the data
- **Filtering**
  - Given the past and the present observation, how should I update my estimate of the true state of nature?
- **Smoothing**
  - Given a complete dataset, what can I infer about the true state of nature in the past?
- **Regression**
  - Given a time series of two phenomena, what is the association between them?

# Neural networks





- Artificial Neural Cells
  - Linear combination of inputs
  - Non-linear activation function
- Connected Neurons
  - Directed graph
  - Layered structure
    - Dense connections
    - Convolutions & pooling
    - Recurrency, signal gates
    - Masking & attention heads
- Universal function approximator
  - Trainable with data
  - Backpropagation
  - Deep vs. shallow architecture

# Neural networks – Input space separation





# Neural networks – Input space separation



$$y = s(\sum w_i x_i - \theta)$$



# Neural networks – Input space separation



$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$




# Neural networks – Input space separation



$$ax + by + c = 0$$

$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$

# Neural networks – Input space separation







$$(a, b) = (w_1, w_2)$$


$$ax + by + c = 0$$

$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$

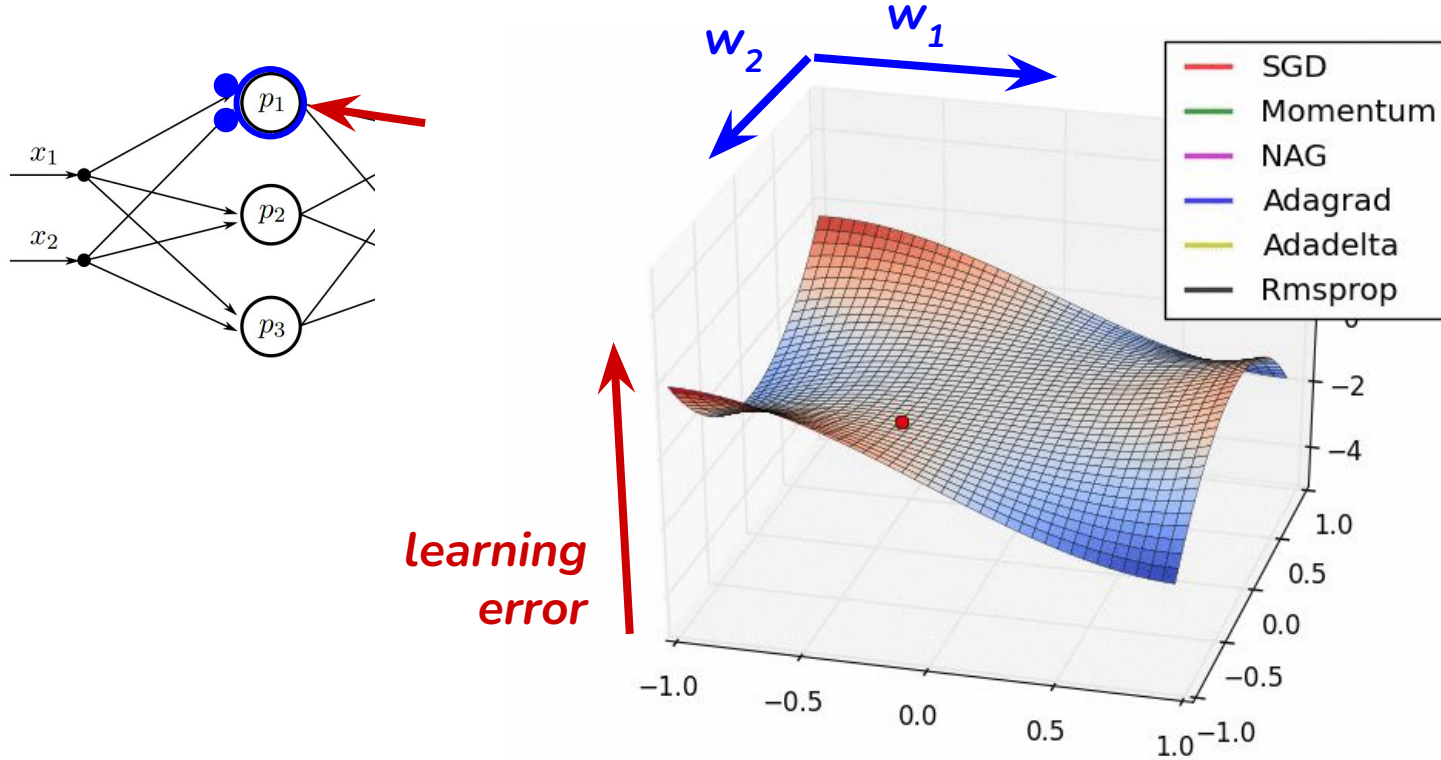
# Neural networks – Input space separation




# Neural networks – Backpropagation



# Neural networks - Backpropagation




# Time Series with Neural Networks




- Neural Networks
  - How to **express time domain**
  - How to **prepare training data**
  - How to **design the model**
  - How to **train & test the model**

# How neural network fits?






# Simple feed-forward network

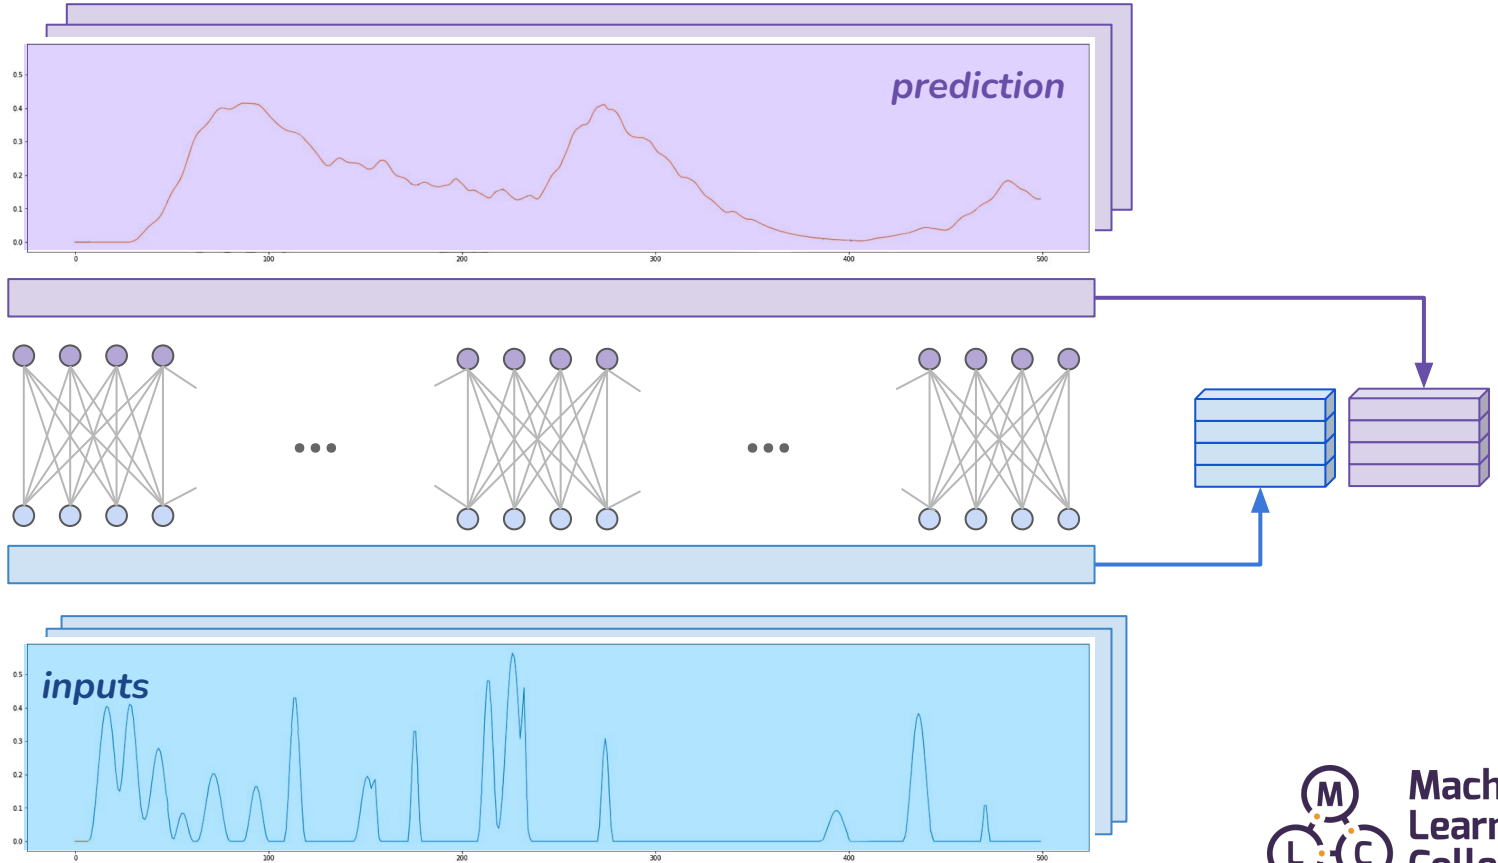





# Simple feed-forward network




# Simple feed-forward network



# Simple feed-forward network




# Simple feed-forward network





*predicting  
from future ?*


# Simple feed-forward network




# Sliding feed-forward network



# Sliding feed-forward network




# Sliding feed-forward network







# Sliding feed-forward network





# Sliding feed-forward network




# Sliding feed-forward network




# Recurrent Neural Networks




# RNN for time series prediction




# RNN for time series prediction



# RNN for time series prediction

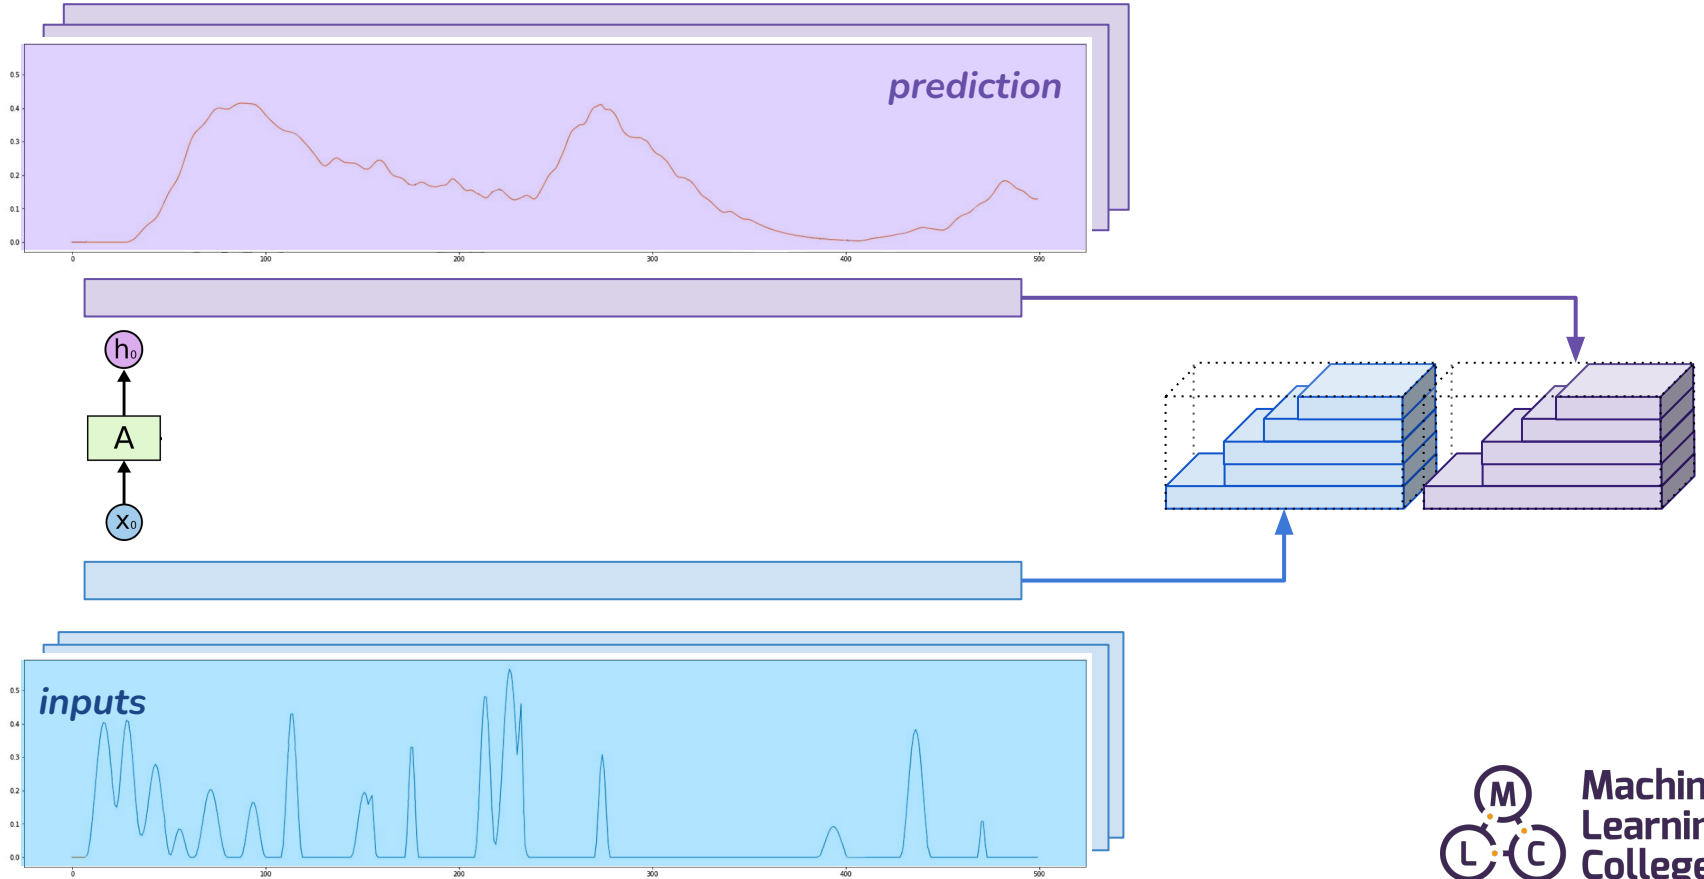


# RNN for time series prediction






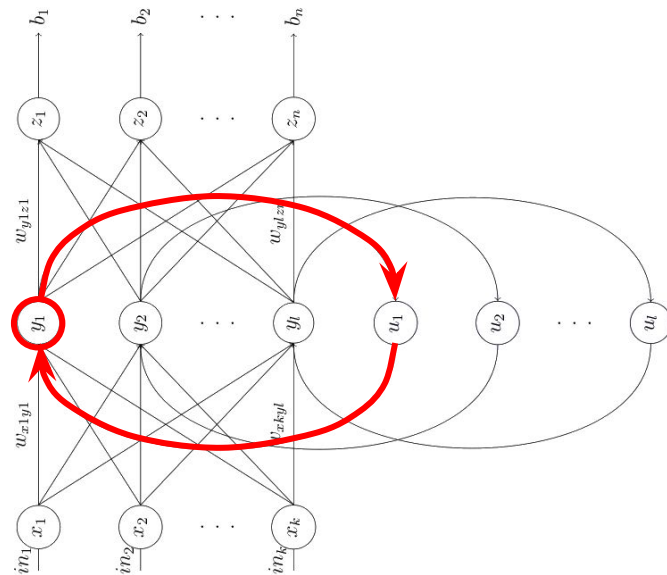
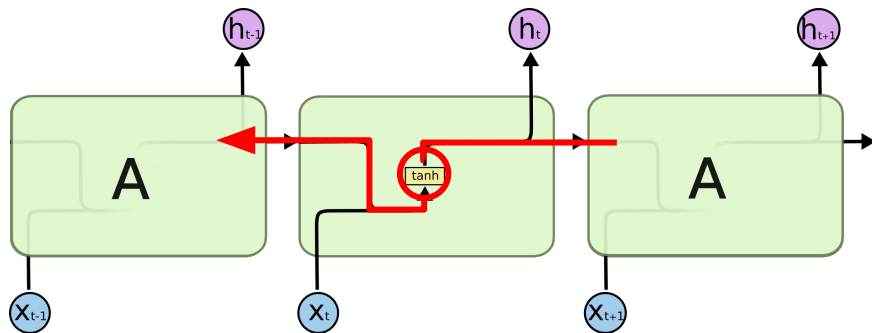
# RNN for time series prediction



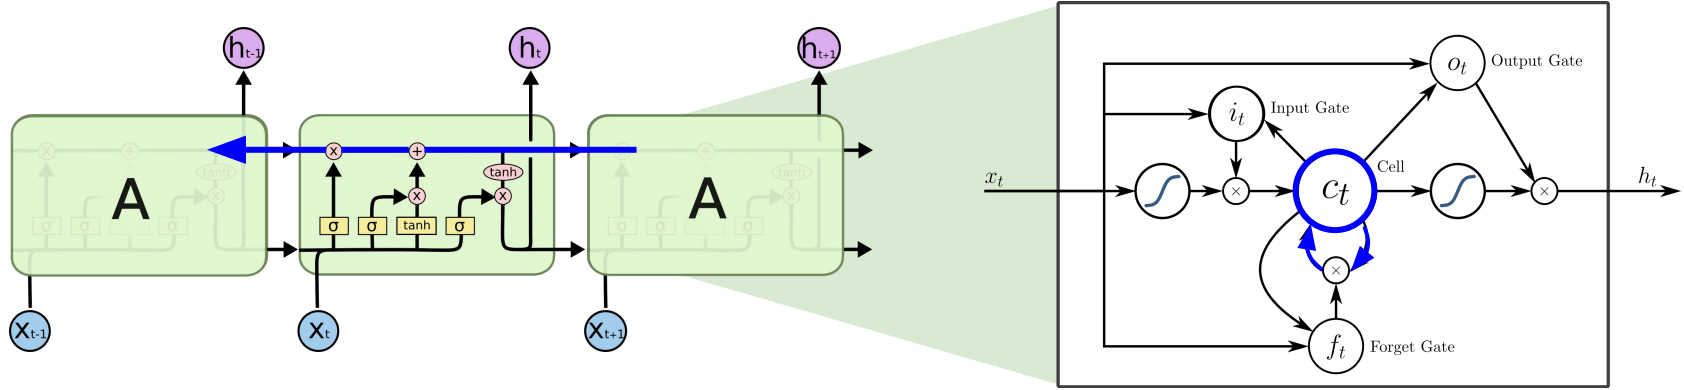
# RNN for time series prediction




# RNN – Vanishing gradients



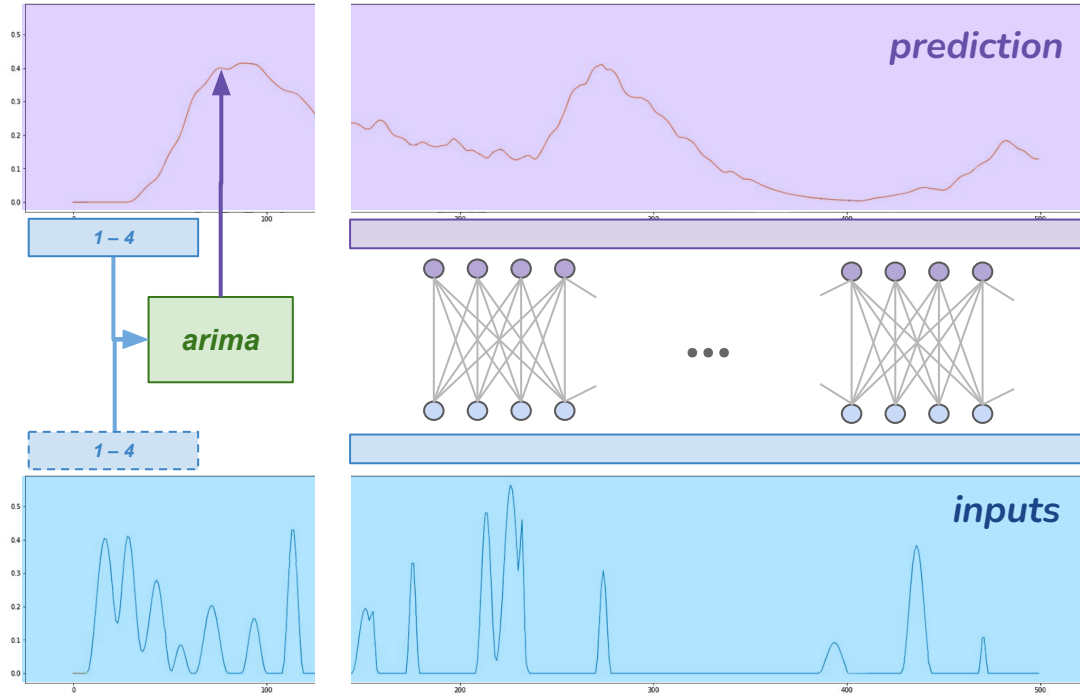
# Long short-term memory – LSTM




# Classical model vs. neural network




# Classical model vs. neural network



# Classical model vs. neural network




# Classical model vs. neural network






# Rainfall-runoff example



- Test out neural networks on simple **generated rainfall-runoff dataset**
  - Simulated **long-time dependencies** in data
  - Test various neural network architectures
    - Flat feed-forward network
    - LSTM

# Tensors & dimensions – univariate regression


*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, T)$$


# Tensors & dimensions – univariate regression

*Flat NN = 2D training data*




# Tensors & dimensions – univariate regression

*Flat NN = 2D training data*



$(B, T) \rightarrow (B, T)$


*Recurrent NN = 3D training data*



$(B, T, 1) \rightarrow (B, T, 1)$


# Tensors & dimensions – univariate regression

*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, T)$$


*Recurrent NN = 3D training data*



$$(B, T, 1) \rightarrow (B, T, 1)$$


# Trampoline jumping example

- Data preparation
  - Dataset normalization
  - Sequence padding
- Binary classification task
  - Target values & dimensions
  - Loss functions
- Training & evaluation
  - Inference visualization
  - Evaluation metrics



# Tensors & dimentions – binary classification


*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, 1)$$


# Tensors & dimensions – **binary classification**

*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, 1)$$

*Recurrent NN = 3D training data*




$$(B, T, 1) \rightarrow (B, 1)$$




# Tensors & dimensions – **binary classification**

*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, 1)$$


*Recurrent NN = 3D training data*



$$(B, T, 1) \rightarrow (B, 1)$$


# Tensors & dimensions – **binary classification**

*Flat NN = 2D training data*



$$(B, T) \rightarrow (B, 1)$$

*Recurrent NN = 3D training data*




$$(B, T, 1) \rightarrow (B, 1)$$

# Tensors & dimentions – multivariate b. classification

Flat NN = 2D training data

?


Recurrent NN = 3D training data




# Tensors & dimensions – multivariate b. classification

Flat NN =  $3D \rightarrow 2D$  training data

flattening / interleaving



Recurrent NN = 3D training data



# Binary classification – confusion matrix

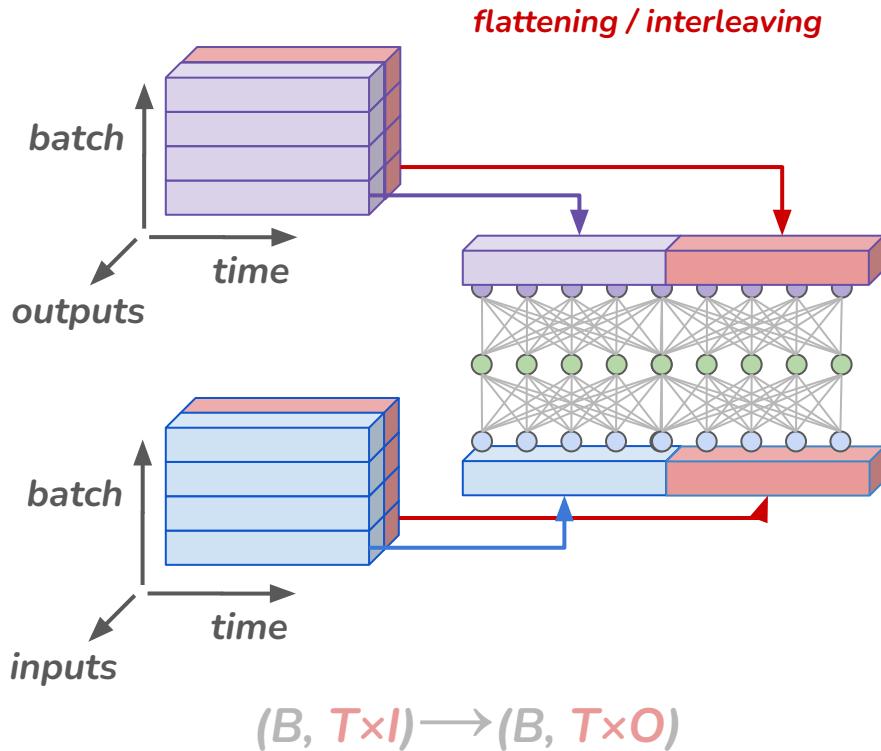
Sources: [20][21][22][23][24][25][26][27] view · talk · edit

		Predicted condition			
		Positive (PP)	Negative (PN)	Informedness, bookmaker informedness (BM) = $TPR + TNR - 1$	Prevalence threshold (PT) = $\frac{\sqrt{TPR \times FPR} - FPR}{TPR - FPR}$
Actual condition	Total population = $P + N$				
	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power = $\frac{TP}{P} = 1 - FNR$	False negative rate (FNR), miss rate = $\frac{FN}{P} = 1 - TPR$
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection	False positive rate (FPR), probability of false alarm, fall-out = $\frac{FP}{N} = 1 - TNR$	True negative rate (TNR), specificity (SPC), selectivity = $\frac{TN}{N} = 1 - FPR$
	Prevalence = $\frac{P}{P + N}$	Positive predictive value (PPV), precision = $\frac{TP}{PP} = 1 - FDR$	False omission rate (FOR) = $\frac{FN}{PN} = 1 - NPV$	Positive likelihood ratio (LR+) = $\frac{TPR}{FPR}$	Negative likelihood ratio (LR-) = $\frac{FNR}{TNR}$
	Accuracy (ACC) = $\frac{TP + TN}{P + N}$	False discovery rate (FDR) = $\frac{FP}{PP} = 1 - PPV$	Negative predictive value (NPV) = $\frac{TN}{PN} = 1 - FOR$	Markedness (MK), deltaP ( $\Delta p$ ) = $PPV + NPV - 1$	Diagnostic odds ratio (DOR) = $\frac{LR+}{LR-}$
	Balanced accuracy (BA) = $\frac{TPR + TNR}{2}$	$F_1$ score = $\frac{2PPV \times TPR}{PPV + TPR} = \frac{2TP}{2TP + FP + FN}$	Fowlkes–Mallows index (FM) = $\sqrt{PPV \times TPR}$	Matthews correlation coefficient (MCC) = $\sqrt{TPR \times TNR \times PPV \times NPV} - \sqrt{FNR \times FPR \times FOR \times FDR}$	Threat score (TS), critical success index (CSI), Jaccard index = $\frac{TP}{TP + FN + FP}$

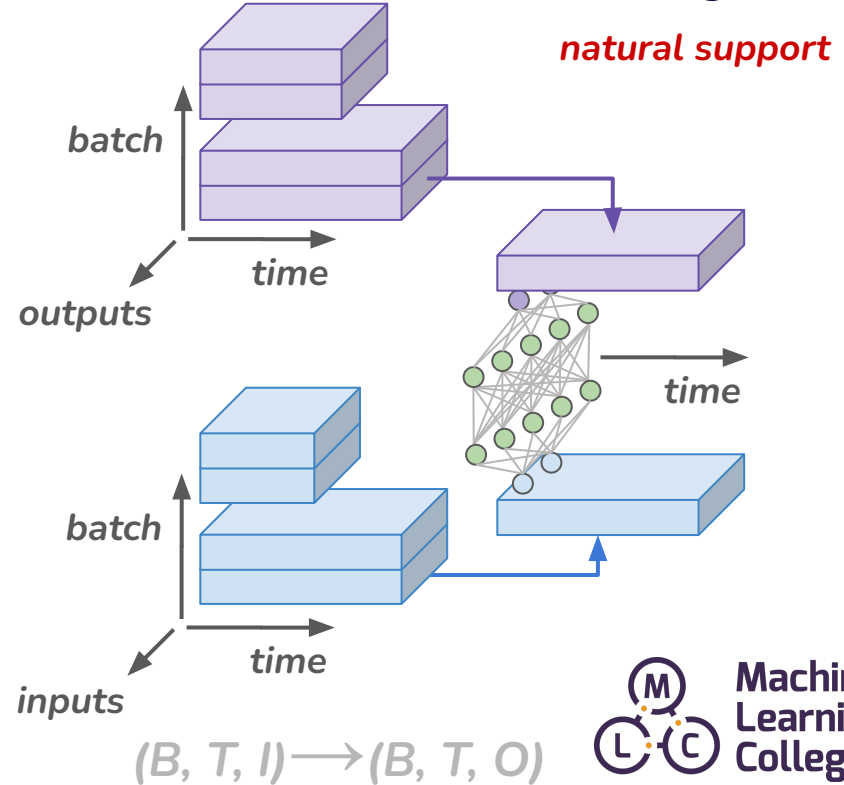
[https://en.wikipedia.org/wiki/Confusion\\_matrix](https://en.wikipedia.org/wiki/Confusion_matrix)

# Tensors & dimensions – multivariate regression

Flat NN =  $3D \rightarrow 2D$  training data



Recurrent NN = 3D training data



# Time Series Modeling using Neural Networks – DAY 2

Dušan Fedorčák  
02/2022 – 02



# Content

## DAY 1

- **Classical time series analysis**
  - *Decomposition of time series*
  - *ARIMA models family*
  - *State space models generalization*
- **Theoretical window**
  - *Neural Networks & Recurrent NNs*
  - *Time series specifics*
- **Practical examples**
  - *Simple regression – toy example*
  - *Rainfall-runoff simulation – regression*

– lunch break –

- **Practical examples**
  - *Trampoline jumping – classification*
  - *Local Weather Forecast – regression*

## DAY 2

- **Product Design & ML**
  - *Integration of ML models into products*
- **Practical Examples** (in random order)
  - *Exoplanets Hunting*
  - *Mobile Motion Sensing*

– lunch break –

- **Tips & tricks for debugging NNs**
- **Practical Examples**
  - *Manufacturing Process Modeling*
  - *Financial distress prediction*



# ML & Product Design – problems & decisions

- ML should solve **problems**
  - AI is cool  $\Rightarrow$  let's use it in the product! ✖
  - We need to solve this problem  $\Rightarrow$  can we apply ML? ✔
  - *Q: What other means of solving the problem are available?*

# ML & Product Design – problems & decisions

- ML should solve **problems**
  - AI is cool  $\Rightarrow$  let's use it in the product! ✖
  - We need to solve this problem  $\Rightarrow$  can we apply ML? ✔
  - *Q: What other means of solving the problem are available?*
- Problems vary in **difficulty**
  - Easily solvable by human  $\Rightarrow$  ML helps to scale up & automate
    - spam filter, face recognition, driving a car
  - Not easily solvable by human  $\Rightarrow$  ML can bring some solution
    - weather forecast, stock market prediction
  - *Q: What are we optimizing for? (costs, risk reduction, better service, ... )*

# ML & Product Design – problems & decisions

- ML should solve **problems**
  - AI is cool  $\Rightarrow$  let's use it in the product! ✖
  - We need to solve this problem  $\Rightarrow$  can we apply ML? ✔
  - *Q: What other means of solving the problem are available?*
- Problems vary in **difficulty**
  - Easily solvable by human  $\Rightarrow$  ML helps to scale up & automate
    - spam filter, face recognition, driving a car
  - Not easily solvable by human  $\Rightarrow$  ML can bring some solution
    - weather forecast, stock market prediction
  - *Q: What are we optimizing for? (costs, risk reduction, better service, ... )*
- Problems boils down to **decisions**
  - ML can **assist** with decisions
  - ML can **automate** decisions
  - *Q: Could assistance model work for us or full automation is needed?*

# ML & Product Design – automation

- ML solutions are **imperfect**
  - Expectation control & automation bias  $\Rightarrow$  trust issues
  - Scaling up imperfect models  $\Rightarrow$  quality issues
  - Right evaluation metrics – model evaluation vs. UX evaluation
  - *Q: Do all involved parties understand the problem & solution?*
  - *Q: Is there an evaluation metric everybody understands and agrees with?*

# ML & Product Design – automation

- ML solutions are **imperfect**
  - Expectation control & automation bias  $\Rightarrow$  trust issues
  - Scaling up imperfect models  $\Rightarrow$  quality issues
  - Right evaluation metrics – model evaluation vs. UX evaluation
  - *Q: Do all involved parties understand the problem & solution?*
  - *Q: Is there an evaluation metric everybody understands and agrees with?*
- **Black-box** models are tricky
  - Black-box model + automation  $\Rightarrow$  trust issues
  - Expensive configurability & finetuning
  - *Q: Can we build an understandable stress-test evaluation dataset?*
  - *Q: What tools or probes are available for analysis of our learned black-box model?*

# ML & Product Design – automation

- ML solutions are **imperfect**
  - Expectation control & automation bias  $\Rightarrow$  trust issues
  - Scaling up imperfect models  $\Rightarrow$  quality issues
  - Right evaluation metrics – model evaluation vs. UX evaluation
  - *Q: Do all involved parties understand the problem & solution?*
  - *Q: Is there an evaluation metric everybody understands and agrees with?*
- **Black-box** models are tricky
  - Black-box model + automation  $\Rightarrow$  trust issues
  - Expensive configurability & finetuning
  - *Q: Can we build an understandable stress-test evaluation dataset?*
  - *Q: What tools or probes are available for analysis of our learned black-box model?*
- Assistance  $\Leftrightarrow$  Automation – there is a **spectrum**
  1. No automation
  2. Scored set of possible decisions
  3. Narrowed set of decision to approve
  4. Veto before automatic execution
  5. Full automation
  - *Q: What is the lowest level of automation that brings value*

# ML & Product Design – data analysis

- ML solutions **depends on data**
  - More complex models  $\Rightarrow$  more data required
  - Constant battle against overfitting
  - Distributions shift over time
  - *Q: What data is available and will be available in the future*

# ML & Product Design – data analysis

- ML solutions **depends on data**
  - More complex models  $\Rightarrow$  more data required
  - Constant battle against overfitting
  - Distributions shift over time
  - *Q: What data is available and will be available in the future*
- **Base rates** in data influences products
  - Base rates can have unintuitive effects on the product
  - Sampling reality often produces imbalanced data
  - *Q: What would be the performance of near-perfect model given the base rates*




# ML & Product Design – data analysis

- ML solutions **depends on data**
  - More complex models  $\Rightarrow$  more data required
  - Constant battle against overfitting
  - Distributions shift over time
  - *Q: What data is available and will be available in the future*
- **Base rates** in data influences products
  - Base rates can have unintuitive effects on the product
  - Sampling reality often produces imbalanced data
  - *Q: What would be the performance of near-perfect model given the base rates*
- All datasets are **biased**
  - Inconsistency between data sampling and model goals
  - Biased evaluation sets
  - *Q: Does our historical data reliably capture the goal of the model*

# Dataset issues – insufficient number of samples



## Dataset issues – insufficient number of samples



# Dataset issues – biased training set

TECHNOLOGY NEWS   OCTOBER 10, 2018 / 5:43 AM / UPDATED 2 YEARS AGO

## Insight - Amazon scraps secret AI recruiting tool that showed bias against women

By Jeffrey Dastin

8 MIN READ



SAN FRANCISCO (Reuters) - Amazon.com Inc's [AMZN.O](#) machine-learning specialists uncovered a big problem: their new recruiting engine did not like women.

<https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G>


# Dataset issues – biased training set



# Dataset issues – biased training set




# Dataset issues – biased training set






# Human perception – overfitted to faces





# Motion sensing example

- Data preparation
  - Dataset normalization
  - Slicing long sequences
- Categorical classification task
  - Predict activity type
  - Use correct activation & loss function
- Training & evaluation
  - Try different architectures
  - Evaluate result with standard metrics
- Secondary task
  - Subject identification



# ML Tips & Tricks

<http://karpathy.github.io/2019/04/25/recipe/>

# ML Tips & Tricks

- **Known your data**
  - Visualize everything you can
  - Try to find patterns  $\Rightarrow$  *become the model yourself*
  - Look for noisy labels / missing data
  - Make sure your preprocessing is correct (especially vectorized code)

# ML Tips & Tricks

- **Known your data**
  - Visualize everything you can
  - Try to find patterns ⇒ *become the model yourself*
  - Look for noisy labels / missing data
  - Make sure your preprocessing is correct (especially vectorized code)
- **Simple models first**
  - Build training & evaluation loop
  - Choose simple architectures first ⇒ *less room for errors*
  - Build baseline models for comparison ⇒ *even simple heuristics are useful*

# ML Tips & Tricks


- **Known your data**
  - Visualize everything you can
  - Try to find patterns  $\Rightarrow$  *become the model yourself*
  - Look for noisy labels / missing data
  - Make sure your preprocessing is correct (especially vectorized code)
- **Simple models first**
  - Build training & evaluation loop
  - Choose simple architectures first  $\Rightarrow$  *less room for errors*
  - Build baseline models for comparison  $\Rightarrow$  *even simple heuristics are useful*
- **Train iteratively**
  - Train without inputs  $\Rightarrow$  *yields another baseline model*
  - Overfit one batch  $\Rightarrow$  *something is off if you can't get zero loss*
  - Overfit the training set as far as you can

# ML Tips & Tricks

- **Known your data**
  - Visualize everything you can
  - Try to find patterns  $\Rightarrow$  *become the model yourself*
  - Look for noisy labels / missing data
  - Make sure your preprocessing is correct (especially vectorized code)
- **Simple models first**
  - Build training & evaluation loop
  - Choose simple architectures first  $\Rightarrow$  *less room for errors*
  - Build baseline models for comparison  $\Rightarrow$  *even simple heuristics are useful*
- **Train iteratively**
  - Train without inputs  $\Rightarrow$  *yields another baseline model*
  - Overfit one batch  $\Rightarrow$  *something is off if you can't get zero loss*
  - Overfit the training set as far as you can
- **Regularize**
  - Early stopping  $\Rightarrow$  *best evaluation loss*
  - Make the model smaller  $\Rightarrow$  *less space for overfitting*
  - Get more training data  $\Rightarrow$  *more labels, data augmentation, pre-training*


# Exoplanets hunting example

- Data preparation
  - Dataset normalization
  - Highly **imbalanced dataset**
- Binary classification/detection task
  - Detect starts with planets
- Model architecture
  - Dense, LSTM, Bidir. LSTM, CNN
- Training & evaluation
  - Use right evaluation metrics for imbalanced datasets



# Factory process example


- Data preparation
  - Dealing with **missing values**
  - Dataset normalization
  - Slicing long sequences
- Regression task
  - Predict target variables in future
- Model architecture
  - **Model architecture mimics the process**
- Training & evaluation
  - Masking out missing labels with **custom loss function**






# RNN and sequence data

one-one




many-one




*sequence  
classification*

one-many




*image  
captioning*

many-many




*time series  
prediction*

many-many




*machine  
translation*


# RNN for time series prediction




# Forecasting from input variables




# Forecasting from input variables




# Forecasting from historical values



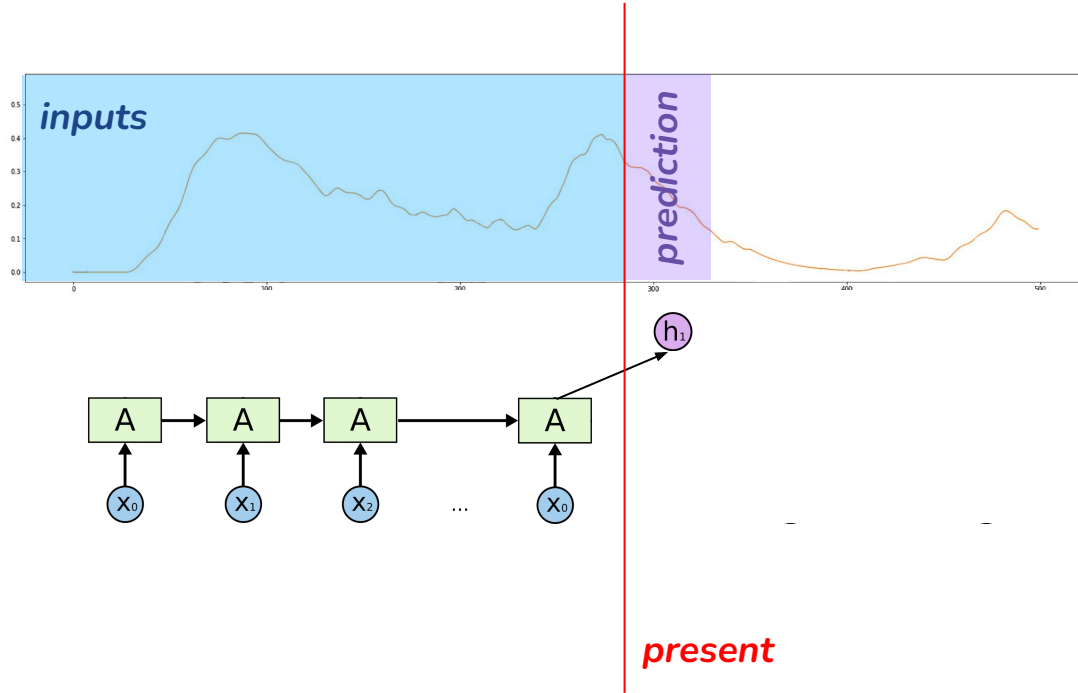
# Forecasting from historical values



# Forecasting from historical values




# Forecasting – one step ahead







# Forecasting – flat multi-step prediction



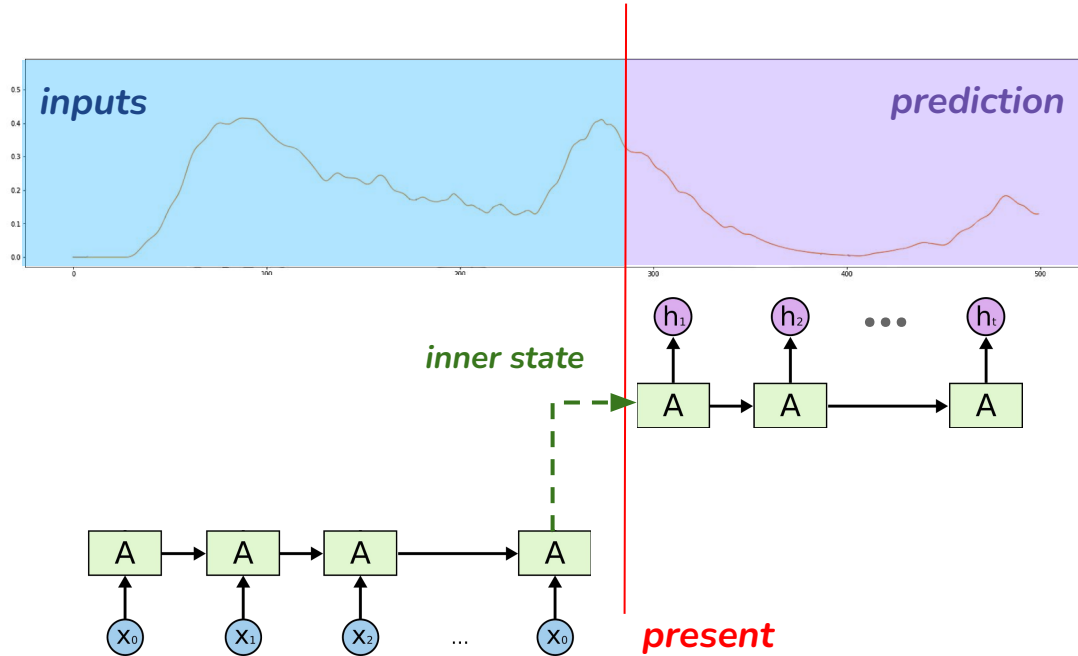
# Forecasting – developed multi-step prediction





# Forecasting – developed multi-step prediction





# Forecasting – encoder & decoder





# Training set construction





# Training set construction



# Training set construction





# Training set construction


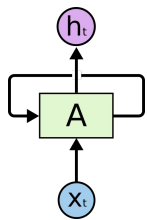







# Training set construction




# Training set construction



# Training set construction




# Training set construction



# Training set construction

training set



testing set



# Training set construction

training set



testing set



# Training set construction

training set


testing set




# Feature Encoding – Seasonal dummy variables

Sun	0	0	0	0
Mon	1	0	0	0
Tue	0	1	0	0
Wed	0	0	1	0
Thu	0	0	0	1
Fri	0	0	0	0
Sat	0	0	0	0

$T \longrightarrow$



- (hour of day, day of week, ...)
- Numerical variables
- One-hot encoding




<https://medium.com/life-at-hopper/ai-in-travel-part-2-representing-cyclic-and-geographic-features-4ada33dd0b22>



[https://kazemnejad.com/blog/transformer\\_architecture\\_positional\\_encoding/](https://kazemnejad.com/blog/transformer_architecture_positional_encoding/)



# Feature Encoding – Seasonal dummy variables



- Circular encoding
- Positional embedding (transformers)

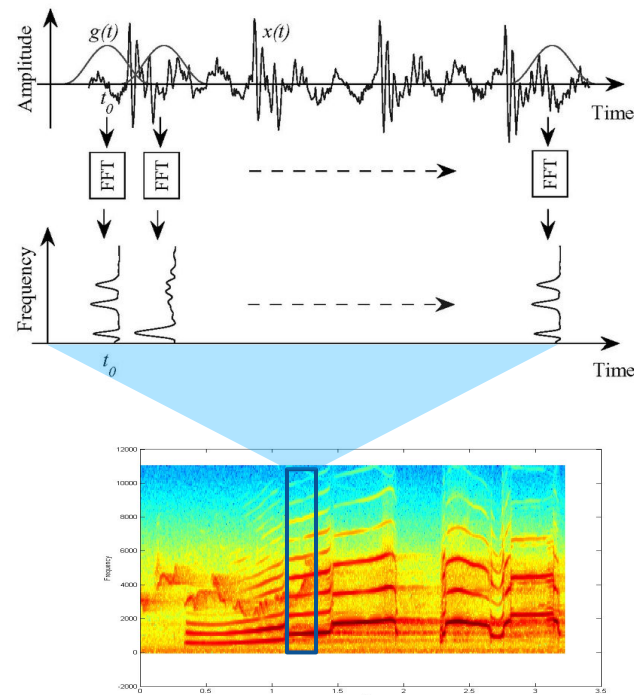



<https://medium.com/life-at-hopper/ai-in-travel-part-2-representing-cyclic-and-geographic-features-4ada33dd0b22>

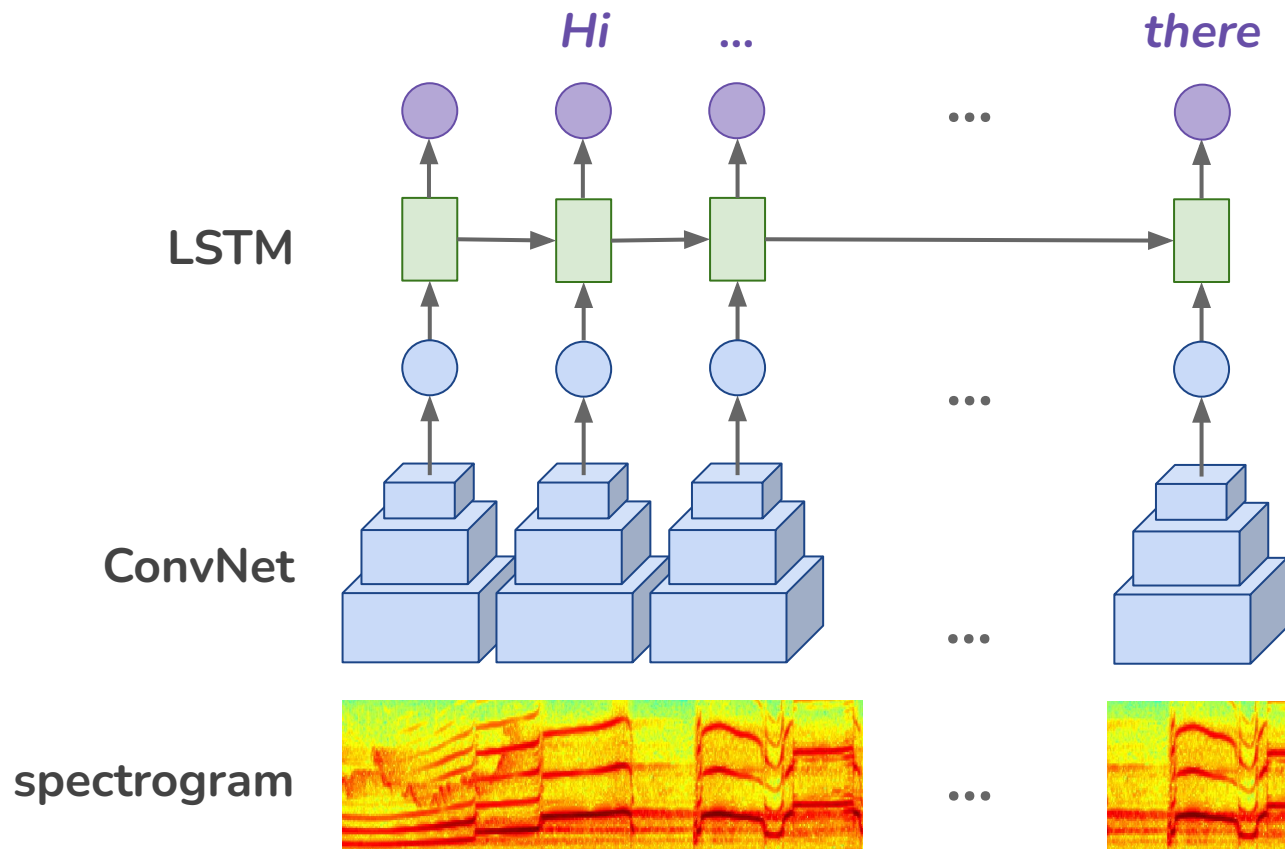
[https://kazemnejad.com/blog/transformer\\_architecture\\_positional\\_encoding/](https://kazemnejad.com/blog/transformer_architecture_positional_encoding/)

# Feature Encoding II

- Exchange for extra dimension
  - 1D Convolution & pooling
  - Short-Time Fourier Transform




# Advanced architectures




# Additional learning materials

- Classical time series
  - [Time series course](#) – a book-like explanation of basic principles of time-series and classical analysis
  - [Statistical forecasting](#) – detailed notes on classical time series analysis and ARIMA models
- Keras
  - [Guides](#) - code examples for most of the basics in Keras
  - [Examples](#) - huge selection of code examples from different areas (time series, vision, ...)
  - [Blog](#) – good selection of advanced application of Keras on practical problems
- Interesting blogs
  - [Adam Geitgey](#) – Machine learning is fun – great selection of simple examples from various areas
  - [Christopher Olah](#) – very well-described principles of neural networks (with a lot of visual insights)
  - [Andrej Karpathy](#) – some very interesting insights (including the debug recipes for NNs)
  - [Distill](#) – Chris Olah and Shan Carter collaboration – open problems in deep learning & advanced topics
- Tech companies blogs
  - [DeepMind \(Google\)](#) – top research in artificial intelligence – usually accompanied with science papers
  - [OpenAI](#) – started as non-commercial research group / answer to Deepmind
  - [Facebook](#) – many interesting projects sometimes with free-to-use pre-learned models
  - [Amazon](#) – many interesting machine learning articles sometimes with detailed papers


# Time series prediction from textual data




# Time series prediction from textual data



# Pre-training with additional data




# Transferring model & exposing feature layer





# Fine-tuning with time series target data



# Fine-tuning with time series target data

