

Time Series Modeling using Neural Networks

Dušan Fedorčák
11/2021 – RBI



Background

- Ph.D. in computer science at VŠB-TU Ostrava
 - Neural networks & unsupervised self-organization
- Experienced in simulations
 - flood prediction system for MSK
 - traffic monitoring & prediction systems
- Experienced in computer graphics & scientific visualization
 - GIS related real-time 3D visualizations
- 5+ years in applied ML and artificial intelligence
 - Lead researcher in GoodAI – general artificial intelligence
 - CTO in Neuron Soundware – sound processing via Deep Learning
 - Lead ML in Merlon Intelligence Inc. – natural language processing

Content

DAY 1

- **Classical time series analysis**
 - *Decomposition of time series*
 - *ARIMA models family*
 - *State space models generalization*
- **Theoretical window**
 - *Neural Networks & Recurrent NNs*
 - *Time series specifics*
- **Practical examples**
 - *Simple regression – toy example*
 - *Rainfall-runoff simulation – regression*

– lunch break –

- **Practical examples**
 - *Trampoline jumping – classification*
 - *Local Weather Forecast – regression*


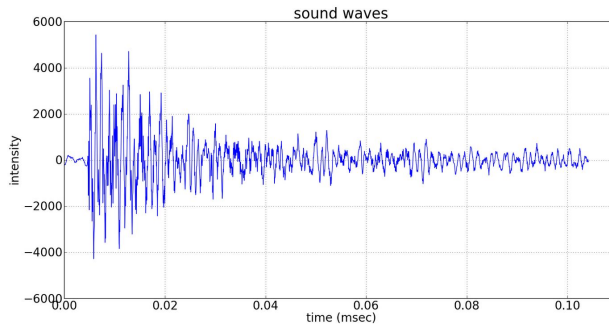

DAY 2

- **Product Design & ML**
 - *Integration of ML models into products*
 - *Tips & tricks for debugging NNs*
- **Practical Examples** (in random order)
 - *Exoplanets Hunting*
 - *Mobile Motion Sensing*

– lunch break –

- *Manufacturing Process Modeling*
- *Financial distress prediction*
- [Google Drive Folder](#) with data
- [GitHub repository](#) with example sources










Time series – example data



CandleStick Chart






Time series – example tasks



Time Series – classical analysis & modeling

- Time Series Decomposition
 - Inflation, trend, seasonality, differencing
- ARIMA models
 - <http://people.duke.edu/~rnau/411home.htm>




State Space Models

- State Space Models
 - A dynamic system that **evolves over time**
 - Knowing **the current state of the model is enough** to predict the future
 - The true state of the system might **not** be **directly observable**
- Model Description
 - State
 - $\mathbf{x}_t \sim N(\mathbf{x}_t, \mathbf{P}_t)$
 - State Equation
 - $\mathbf{x}_t = \mathbf{F}\mathbf{x}_{t-1} + N(0, \mathbf{Q})$ – sometimes without noise
 - Observation Equation
 - $\mathbf{y}_t = \mathbf{H}\mathbf{x}_t + N(0, \mathbf{R})$

Kalman Filter


- Evolve state
 - $\mathbf{x}'_t = F\mathbf{x}_{t-1}$
 - $\mathbf{P}'_t = F\mathbf{P}_{t-1}F^T + Q$
- Integrate observation
 - $\mathbf{x}_t = \mathbf{x}'_t + K_t(y_t - H\mathbf{x}'_t)$
 - $\mathbf{P}_t = (I - K_tH)\mathbf{P}'_t$
- Kalman Gain
 - $K_t = \mathbf{P}'_t H^T (H\mathbf{P}'_t H^T + R)^{-1}$
- ARIMA and Kalman Filter
 - ARIMA can be viewed as a state space model
 - ARIMA can be fitted with MLE via Kalman Filter
 - <https://bookdown.org/rdpeng/timeseriesbook/maximum-likelihood-with-the-kalman-filter.html>
 - <https://towardsdatascience.com/the-kalman-filter-and-maximum-likelihood-9861666f6742>



Hidden Markov Model

- Model Description

- HMM (λ) can be viewed as a **state space model**
- Finite set of hidden states
 - $Q = \{q_1, q_2, \dots, q_n\}$, $\pi = \{\pi_1, \pi_2, \dots, \pi_n\}$ – init
 - n – number of states (hyperparameter)
- Set of observations
 - $O_i = (o^1, o^2, o^3, \dots, o^T)$
- Transition probability matrix & emissions
 - $A = (a_{00}, \dots, a_{nn})$, $B = q_i \rightarrow o$



- Model Capabilities

- $P(O|\lambda)$ – Give prob. of O being produced by λ – *forward-backward alg.*
- $P(q_1, \dots, q_t|O, \lambda)$ – Give most likely sequence of states for given O – *Viterbi alg.*
- $O \Rightarrow \lambda$ – Model must be trainable with O – *Baum-Welch alg.*


Kalman Filter vs. Hidden Markov Model

- **Kalman Filter**

- Continuous state
- Generic state & observation equation
- Linear dynamic system
- *Fusion of sensor readings and controls*
- *ARMA models implementation*

- **Hidden Markov Model**

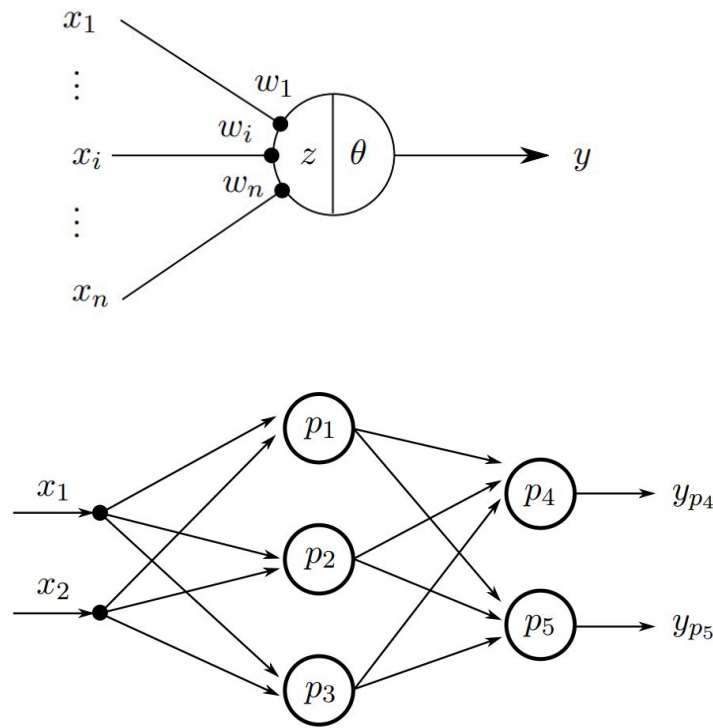
- Discrete set of states
- N-states hyperparameter
- Emission & Transition tables
- *Speech recognition*
- *Time series segmentation*



Time Series – goal in classical terminology



- **Forecasting**
 - Given the past and the present observation, what will the future look like?
- **Time scale analysis**
 - Given the observations, what time scales dominate when observing temporal variation in the data
- **Filtering**
 - Given the past and the present observation, how should I update my estimate of the true state of nature?
- **Smoothing**
 - Given a complete dataset, what can I infer about the true state of nature in the past?
- **Regression**
 - Given a time series of two phenomena, what is the association between them?

Neural networks





- Artificial Neural Cells
 - Linear combination of inputs
 - Non-linear activation function
- Connected Neurons
 - Directed graph
 - Layered structure
 - Dense connections
 - Convolutions & pooling
 - Recurrency, signal gates
 - Masking & attention heads
- Universal function approximator
 - Trainable with data
 - Backpropagation
 - Deep vs. shallow architecture

Neural networks – Input space separation





Neural networks – Input space separation





$$y = s(\sum w_i x_i - \theta)$$

Neural networks – Input space separation



$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$


Neural networks – Input space separation



$$ax + by + c = 0$$

$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$

Neural networks – Input space separation







$$(a, b) = (w_1, w_2)$$


$$ax + by + c = 0$$

$$y = s(\sum w_i x_i - \theta) = s(w_1 x_1 + w_2 x_2 - \theta)$$

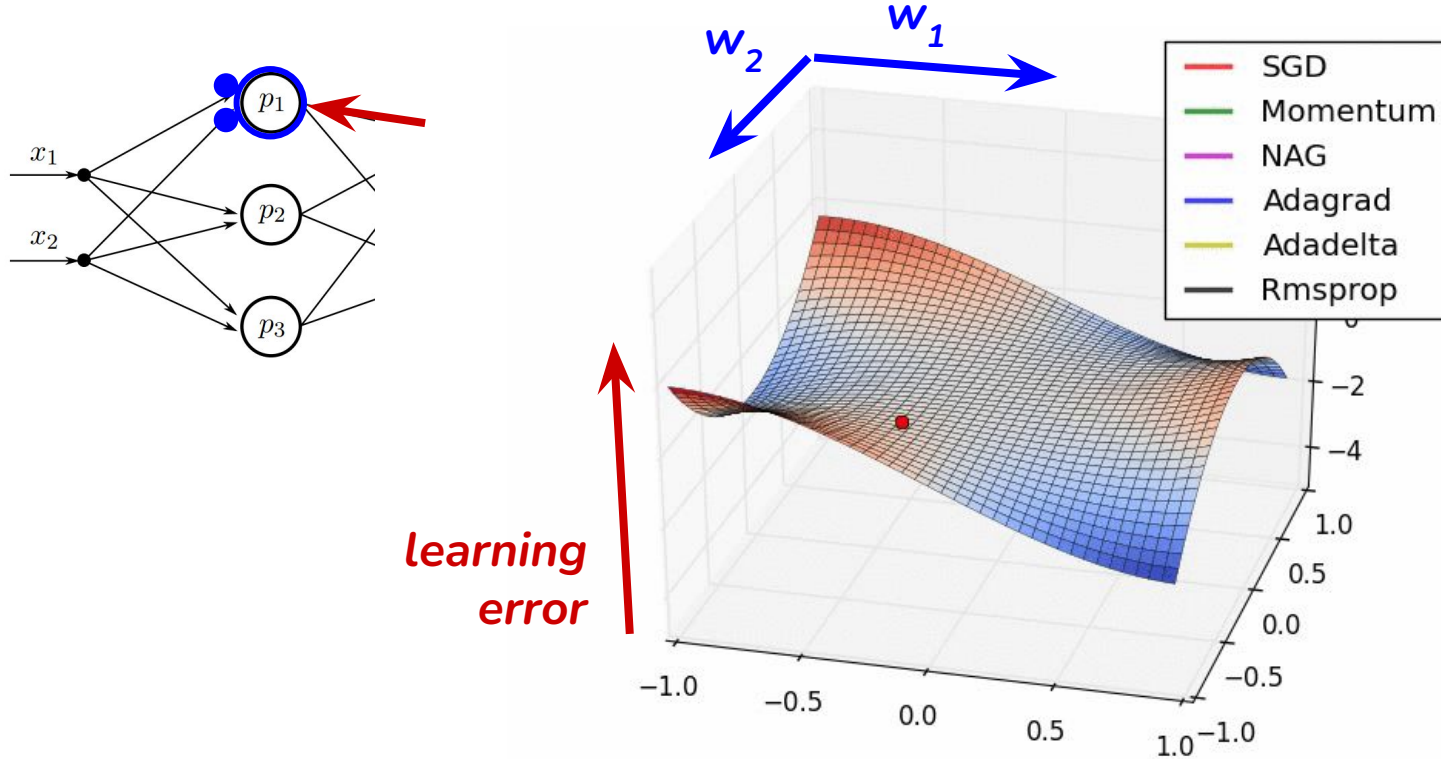
Neural networks – Input space separation




Neural networks – Backpropagation



Neural networks - Backpropagation




Time Series with Neural Networks




- Neural Networks
 - How to **express time domain**
 - How to **prepare training data**
 - How to **design the model**
 - How to **train & test the model**




How neural network fits?



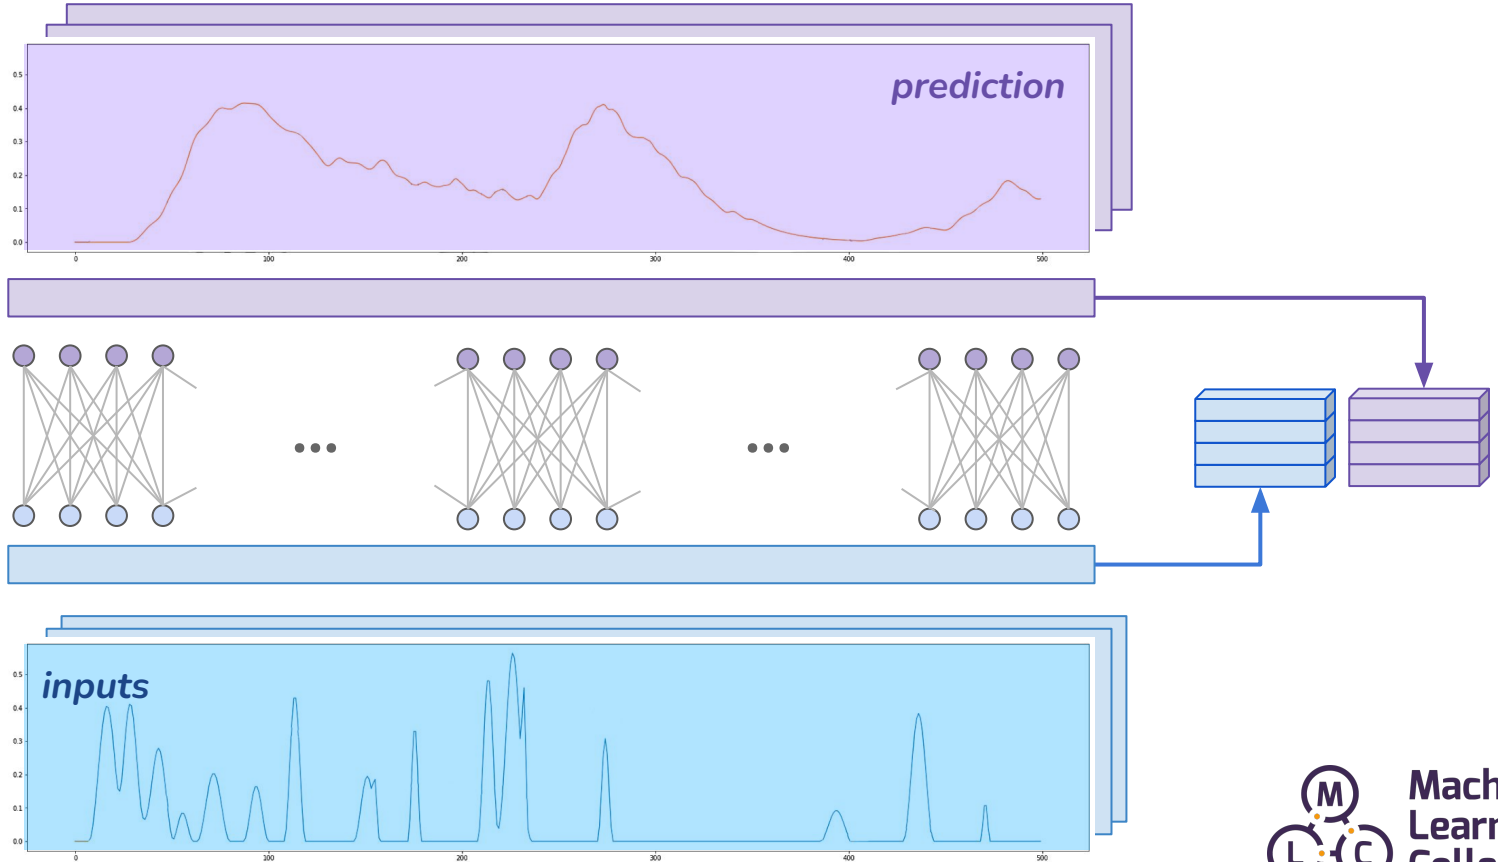
Simple feed-forward network




Simple feed-forward network




Simple feed-forward network



Simple feed-forward network




Simple feed-forward network






*predicting
from future ?*


Simple feed-forward network




Sliding feed-forward network




Sliding feed-forward network




Sliding feed-forward network




Sliding feed-forward network







Sliding feed-forward network




Sliding feed-forward network




Recurrent Neural Networks




RNN for time series prediction




RNN for time series prediction




RNN for time series prediction




RNN for time series prediction





RNN for time series prediction




RNN for time series prediction



RNN – Vanishing gradients




Long short-term memory – LSTM



Rainfall-runoff example


rainfall_runoff.ipynb



- Test out neural networks on simple **generated rainfall-runoff dataset**
 - Simulated **long-time dependencies** in data
 - Test various neural network architectures
 - Flat feed-forward network
 - LSTM


Trampoline jumping example

- Data preparation
 - Dataset normalization
 - Sequence padding
- Binary classification task
 - Target values & dimensions
 - Loss functions
- Training & evaluation
 - Inference visualization
 - Evaluation metrics



Tensors & dimensions – univariate regression


Flat NN = 2D training data



$$(B, T) \rightarrow (B, T)$$


Tensors & dimensions – univariate regression

Flat NN = 2D training data




Tensors & dimentions – univariate regression

Flat NN = 2D training data



$$(B, T) \rightarrow (B, T)$$


Recurrent NN = 3D training data



$$(B, T, 1) \rightarrow (B, T, 1)$$


Tensors & dimensions – univariate regression

Flat NN = 2D training data



$$(B, T) \rightarrow (B, T)$$

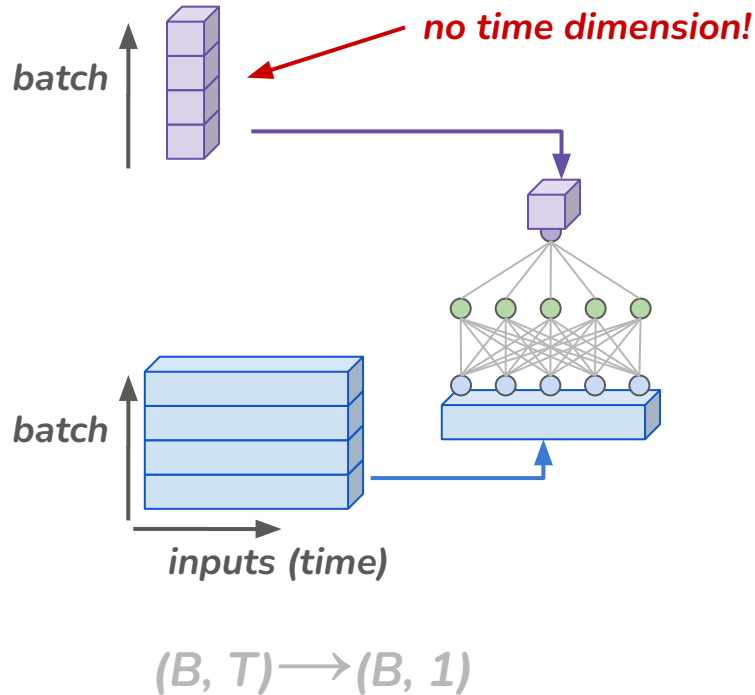
Recurrent NN = 3D training data



$$(B, T, 1) \rightarrow (B, T, 1)$$


Tensors & dimentions – binary classification

Flat NN = 2D training data




Tensors & dimensions – **binary classification**

Flat NN = 2D training data



$$(B, T) \rightarrow (B, 1)$$


Recurrent NN = 3D training data



$$(B, T, 1) \rightarrow (B, 1)$$


Tensors & dimensions – binary classification

Flat NN = 2D training data



$$(B, T) \rightarrow (B, 1)$$


Recurrent NN = 3D training data



$$(B, T, 1) \rightarrow (B, 1)$$


Tensors & dimensions – **binary classification**

Flat NN = 2D training data



$$(B, T) \rightarrow (B, 1)$$

Recurrent NN = 3D training data




$$(B, T, 1) \rightarrow (B, 1)$$

Tensors & dimentions – multivariate b. classification

Flat NN = 2D training data

?


Recurrent NN = 3D training data



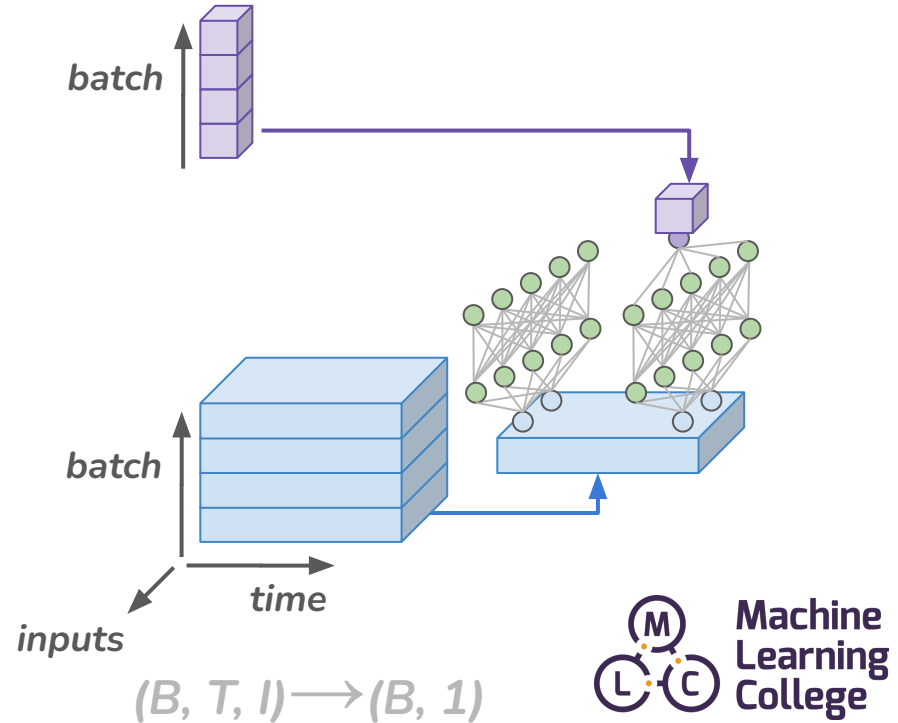
Tensors & dimensions – multivariate b. classification

Flat NN = $3D \rightarrow 2D$ training data

flattening / interleaving




Recurrent NN = 3D training data

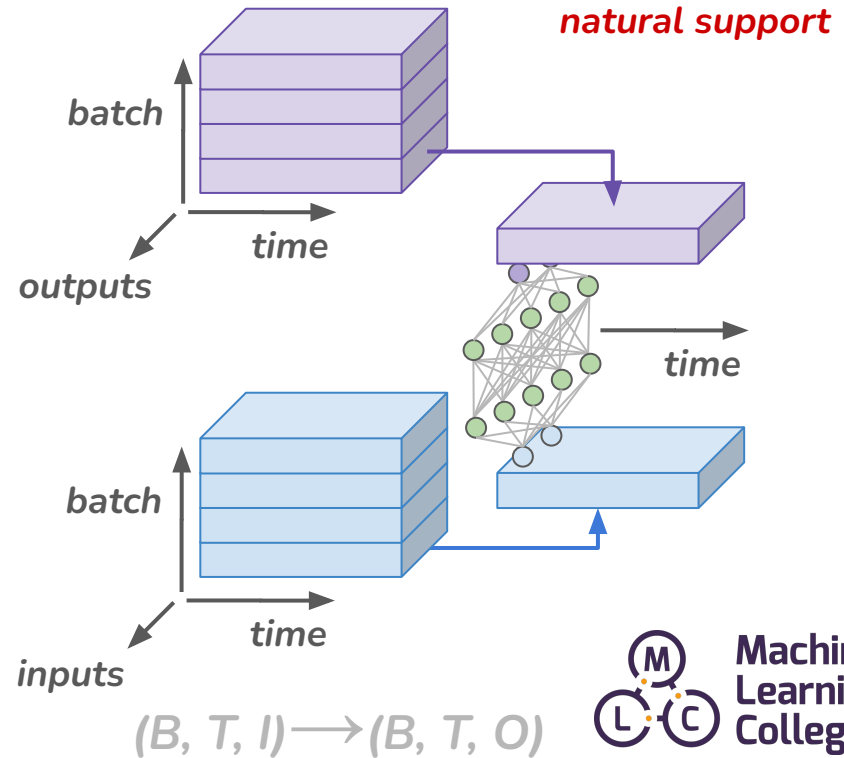


Tensors & dimensions – multivariate regression

Flat NN = $3D \rightarrow 2D$ training data




Recurrent NN = 3D training data




RNN and sequence data

one-one




many-one




*sequence
classification*

one-many




*image
captioning*

many-many



*time series
prediction*


many-many




*machine
translation*

Motion sensing example


- Data preparation
 - Dataset normalization
 - Slicing long sequences
- Categorical classification task
 - Predict activity type
 - Use correct activation & loss function
- Training & evaluation
 - Try different architectures
 - Evaluate result with standard metrics
- Secondary task
 - Subject identification




Forecasting




Forecasting




Forecasting




Forecasting




Forecasting




Forecasting




Forecasting



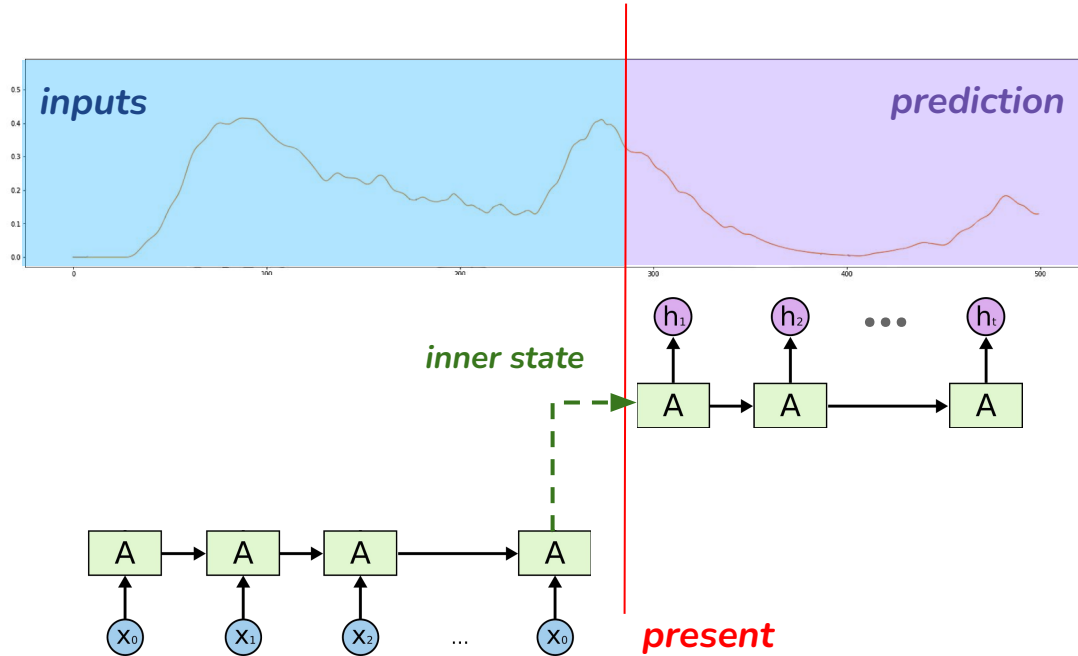
Forecasting





Forecasting





Forecasting – encoder & decoder





Training set construction





Training set construction





Training set construction




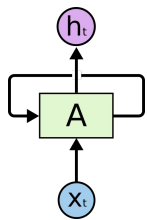
Training set construction






Training set construction




Training set construction



Training set construction




Training set construction



Training set construction

training set



testing set



Training set construction

training set



testing set



Training set construction

training set


testing set




Feature Encoding – Seasonal dummy variables

Sun	0	0	0	0
Mon	1	0	0	0
Tue	0	1	0	0
Wed	0	0	1	0
Thu	0	0	0	1
Fri	0	0	0	0
Sat	0	0	0	0

$T \longrightarrow$




- (hour of day, day of week, ...)
- Numerical variables
- One-hot encoding





<https://medium.com/life-at-hopper/ai-in-travel-part-2-representing-cyclic-and-geographic-features-4ada33dd0b22>

https://kazemnejad.com/blog/transformer_architecture_positional_encoding/

Feature Encoding – Seasonal dummy variables



- Circular encoding
- Positional embedding (transformers)

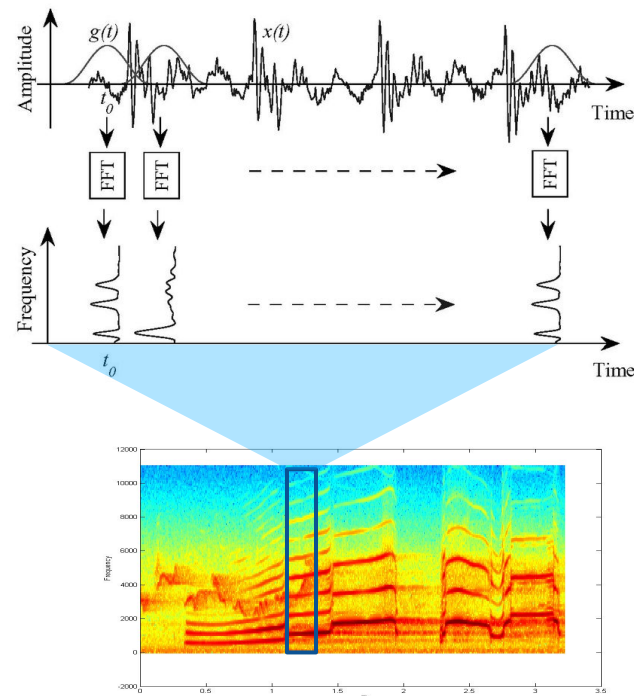



<https://medium.com/life-at-hopper/ai-in-travel-part-2-representing-cyclic-and-geographic-features-4ada33dd0b22>


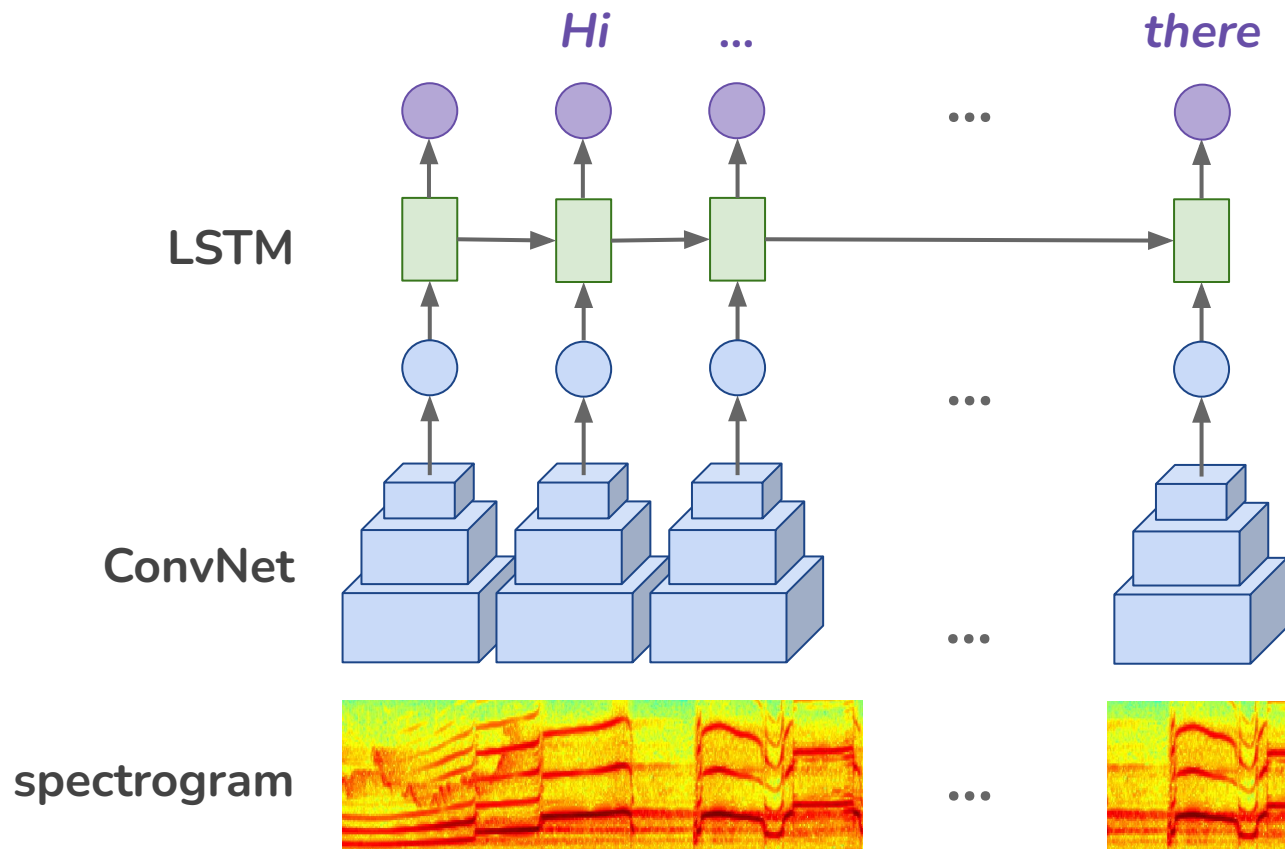
https://kazemnejad.com/blog/transformer_architecture_positional_encoding/

Feature Encoding II


- Exchange for extra dimension
 - 1D Convolution & pooling
 - Short-Time Fourier Transform




Speech recognition




Time series prediction from textual data



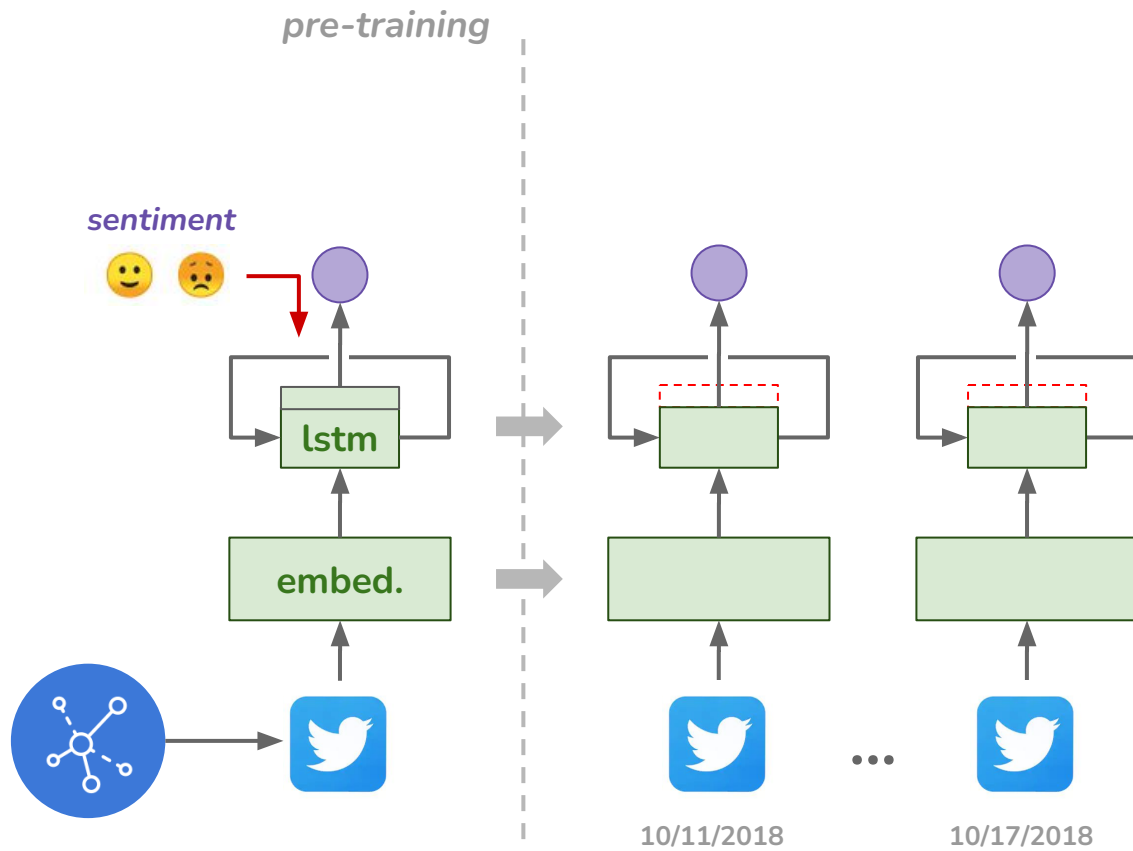
Time series prediction from textual data



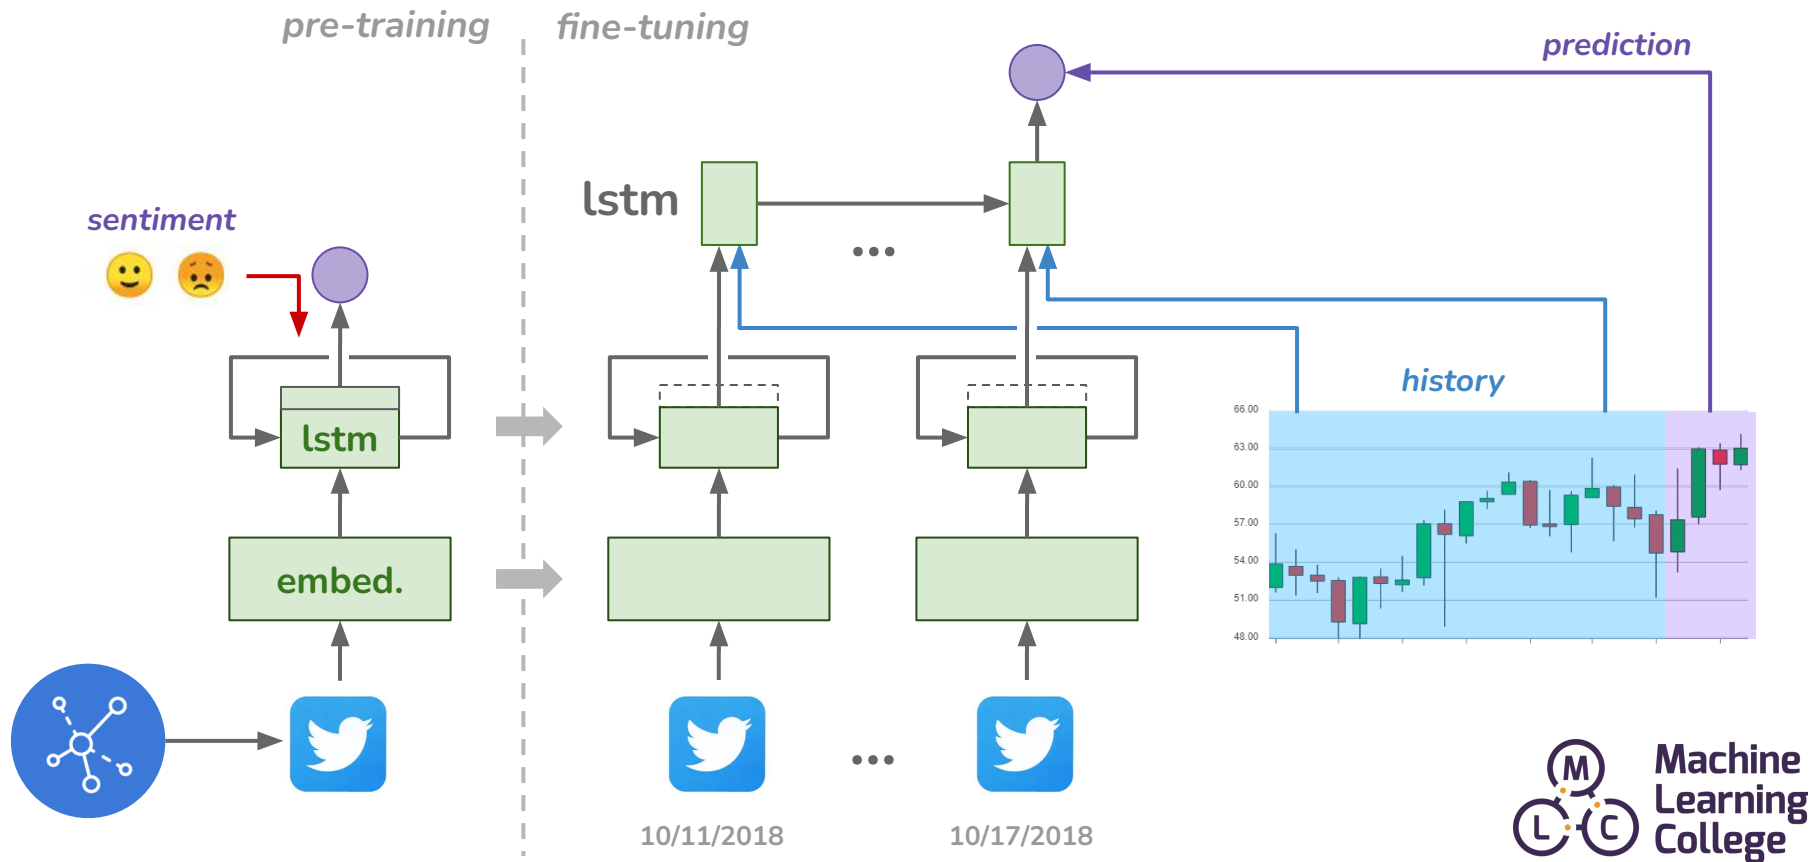
Pre-training with additional data



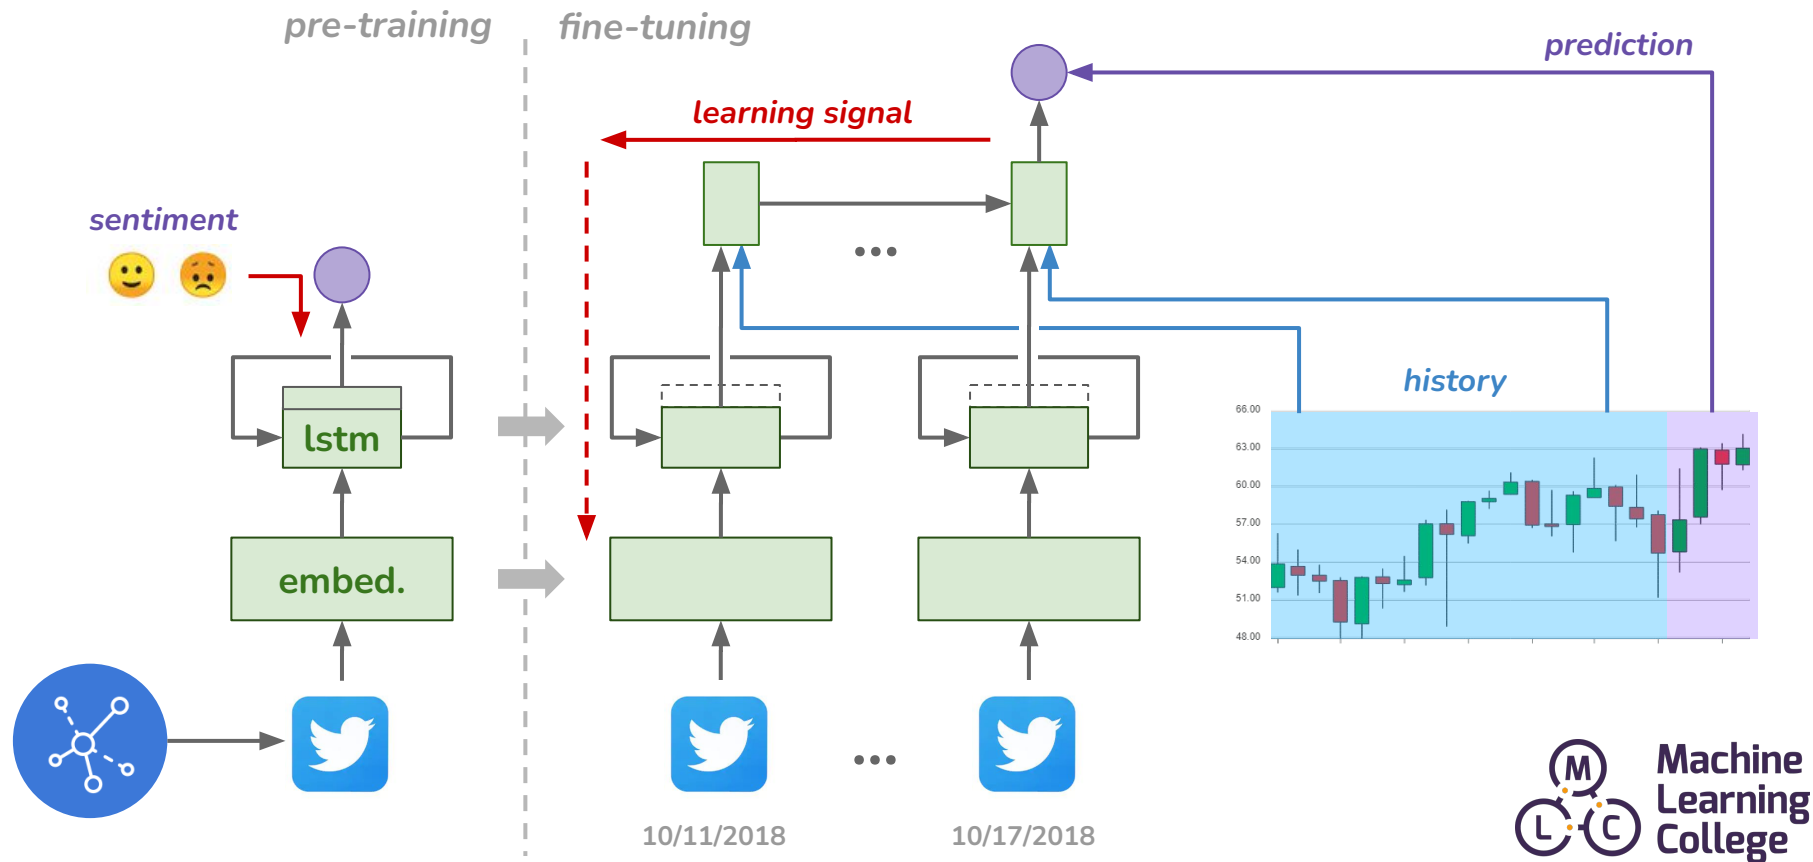
Transferring model & exposing feature layer



Fine-tuning with time series target data



Fine-tuning with time series target data



ML & Product Design

- ML should solve problems
 - *Easily solvable by human* \Rightarrow scale up & automate
 - *Not easily solvable by human* \Rightarrow too complex or not sensitive enough
- Automation is a spectrum
 - Collaboration instead of full automation
 - No automation
 - Scored set of possible decisions
 - Narrowed set of decision to approve
 - Veto before automatic execution
 - Full automation
- Models are imperfect
 - Right evaluation metrics
 - Expectation control / Automation bias
- End-to-end models?
 - Explainability
 - Configurability

ML Tips & Tricks

- Known your data
 - Visualize everything you can
 - Try to find patterns \Rightarrow *become the model yourself*
 - Look for noisy labels / missing data
 - Make sure your preprocessing is correct (especially vectorized code)
- Start with simple models
 - Build training & evaluation loop
 - Choose simple architectures first \Rightarrow *less room for errors*
 - Build baseline models for comparison \Rightarrow *even simple heuristics are useful*
- Train iteratively
 - Train without inputs \Rightarrow *yields another baseline model*
 - Overfit one batch \Rightarrow *something is off if you can't get zero loss*
 - Overfit the training set as far as you can
- Regularize
 - Early stopping \Rightarrow *best evaluation loss*
 - Make the model smaller \Rightarrow *less space for overfitting*
 - Get more training data \Rightarrow *more labels, data augmentation, pre-training*